

Lecture 13: **Graph ADT and Theory**

01204212 Abstract Data Types and Problem Solving

Department of Computer Engineering
Faculty of Engineering, Kasetsart University
Bangkok, Thailand.




Department of
Computer Engineering
Kasetsart University



Outline

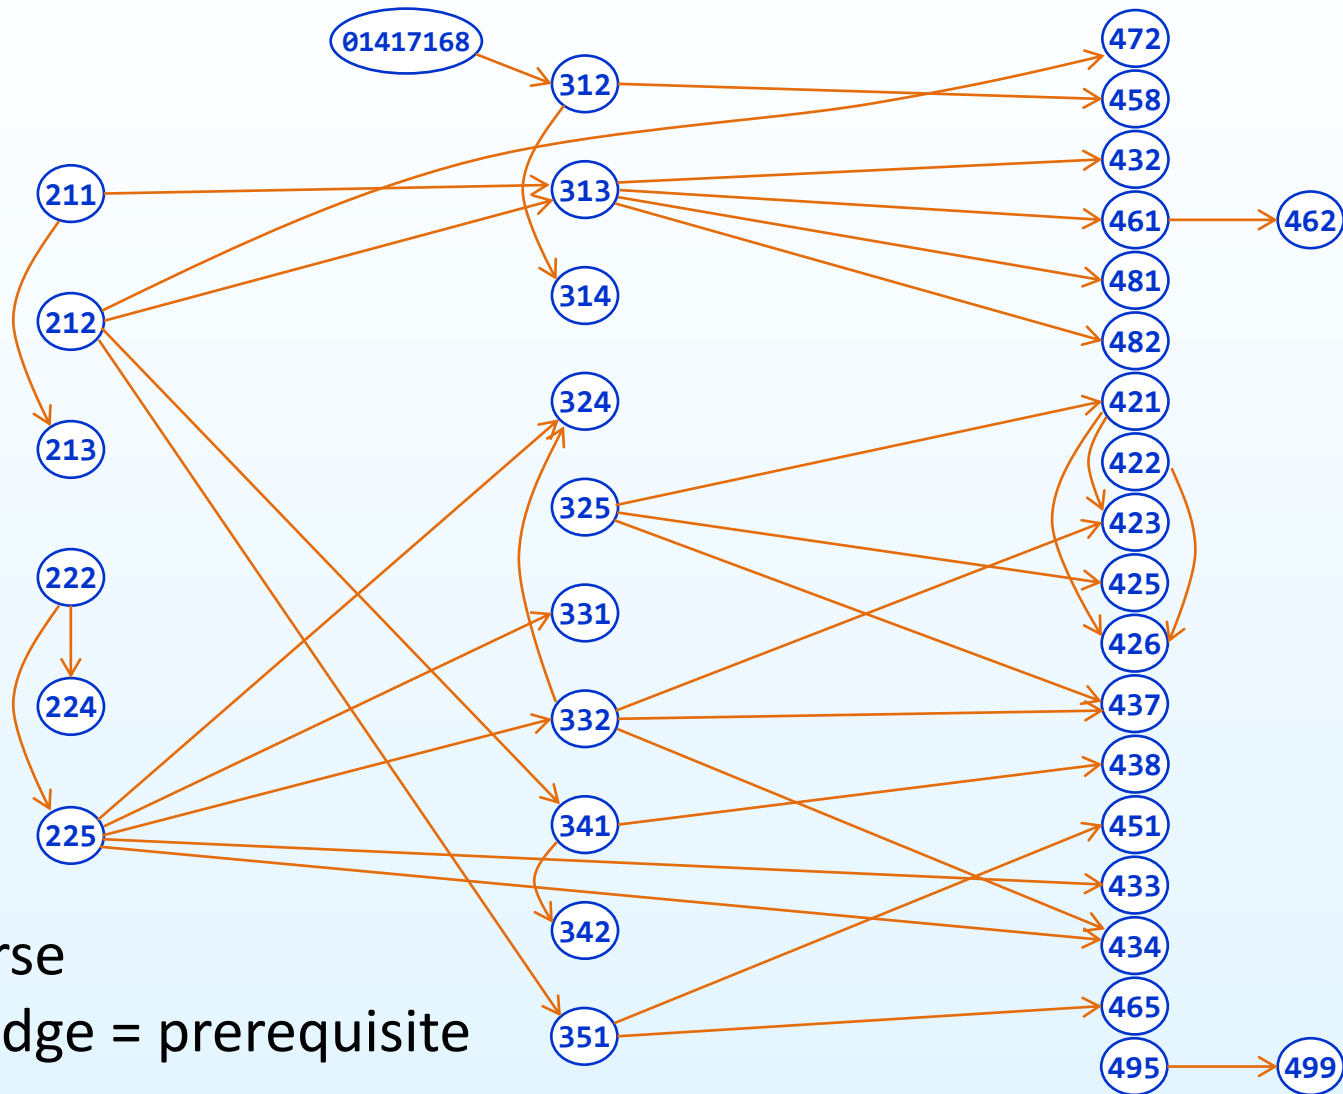
- Graph ADT and Representations
- Graph Traversals
- Connectivity
- Bipartite Graphs
- Topological Sort

What is a Graph?





- Yes, this is a graph
- But we are interested in a different kind of graph
- Consider graphs representing the following problems ...

CPE Course Prerequisites



node = course
(directed) edge = prerequisite


Maze Representation



node = room

edge = door or passage


Electrical Circuit



node = battery, switch, resistor, etc.
edge = connection

Precedence

```
a = 0;      // s1
b = 1;      // s2
c = a+1;    // s3
d = b+a;    // s4
e = d+1;    // s5
e = c+d;    // s6
```



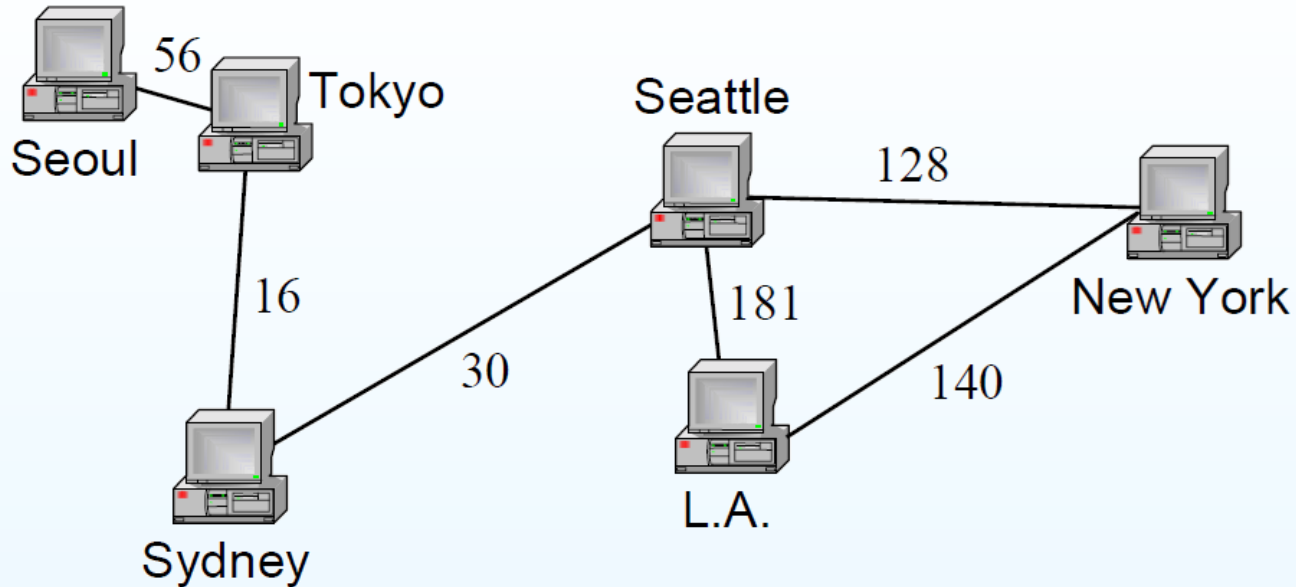
Which statements must execute before s6?

➡ s1, s2, s3, s4

node = statement

edge = precedence requirement

Information Transmission in a Computer Network



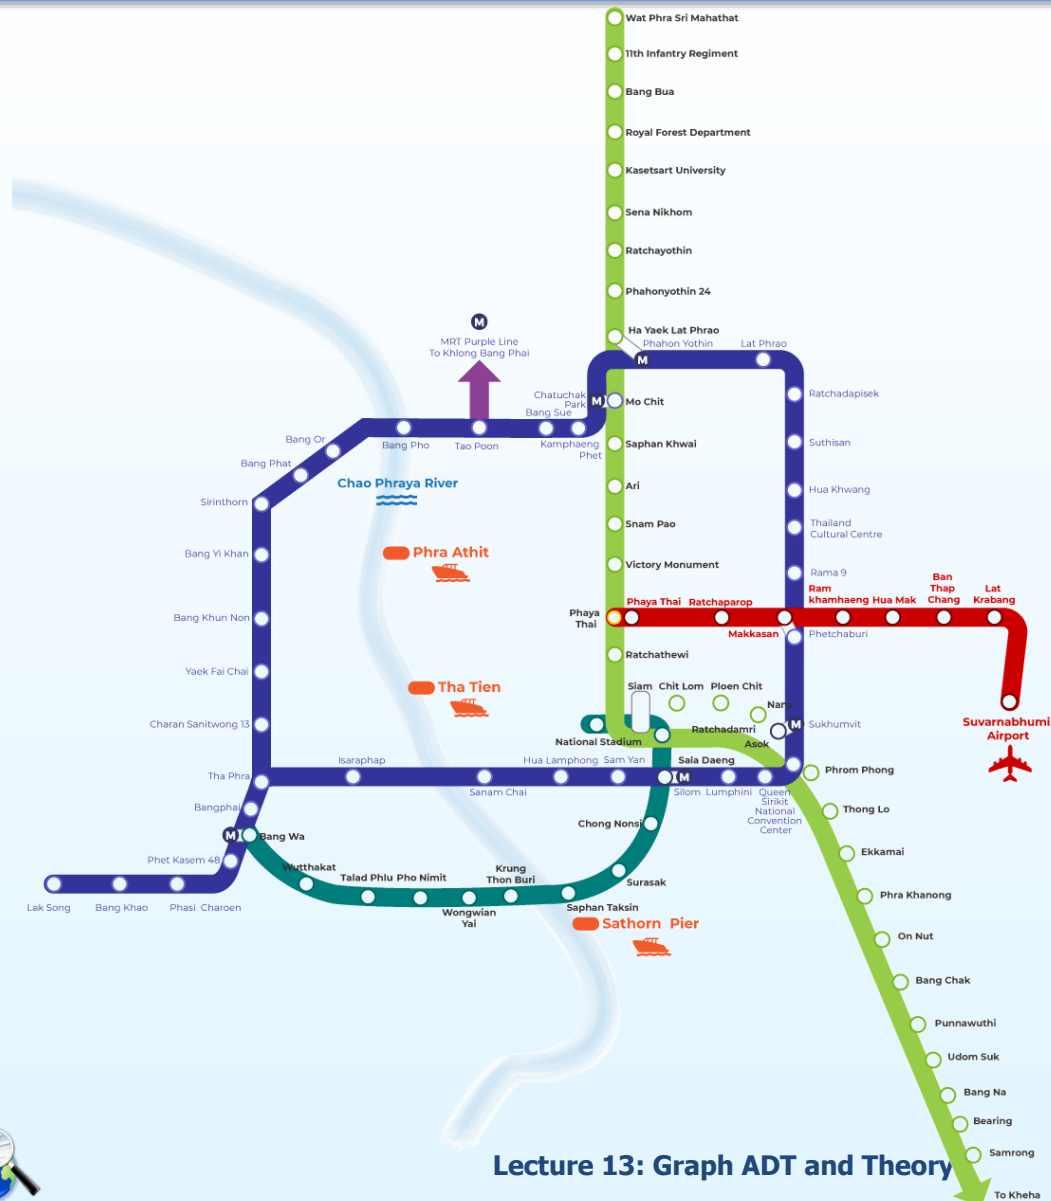
node = computer

edge = connection


weight = transmission rate

Bangkok BTS and MRT Map

node = station
edge = linkage



Molecules



node = atom
edge = bond

Graph Definition


- A **graph** is a collection of **nodes** and **edges**
- Formally,
 - A graph $G = (V, E)$ where
 - V is a set of vertices or nodes
 - $E \subseteq V \times V$ is a set of edges that connect vertices

Graph ADT

- **Graph ADT** describes a container storing an adjacency relation
- **Operations** include:
 - Inserting or removing a vertex (and all edges connecting that vertex)
 - Inserting or removing an edge
 - Querying
 - The number of vertices
 - The number of edges
 - The vertices adjacent to a given vertex
 - The question that “are two vertices adjacent?”
 - The question that “are two vertices connected?”
- The running time of these operations will depend on the representation


Graph Example

- Consider a collection of vertices $V = \{v_1, v_2, \dots, v_9\}$ where $|V| = 9$
- Associated with these vertices are $|E| = 5$ edges
 $E = \{(v_1, v_2), (v_3, v_5), (v_4, v_8), (v_4, v_9), (v_6, v_9)\}$
 - The pair (v_i, v_j) indicates that two vertices v_i and v_j are adjacent




Undirected and Directed Graphs

- If the order of edge pairs (v_i, v_j) does not matter, the graph is an **undirected graph**: $(v_i, v_j) = (v_j, v_i)$



- If the order of edge pairs (v_i, v_j) matters, the graph is a **directed graph** (or called a **digraph**): $(v_i, v_j) \neq (v_j, v_i)$



Maximum Number of Edges

- For an **undirected** graph, the maximum number of edges is

$$|E| \leq \binom{|V|}{2} = \frac{|V|(|V| - 1)}{2} = O(|V|^2)$$



- For a **directed** graph, the maximum number of edges is

$$|E| \leq 2 \binom{|V|}{2} = 2 \frac{|V|(|V| - 1)}{2} = O(|V|^2)$$

Graph Example

Given $G = (V, E)$ where $V = \{A, B, C, D, E, F, G\}$ and $E = \{(A, B), (A, D), (A, E), (B, C), (B, D), (B, E), (C, E), (C, F), (D, E)\}$


- If G is an **undirected** graph
- If G is a **directed** graph



Degrees

- If G is an **undirected** graph

Degree: # of edges incident with



$$\deg(A) = \deg(C) = \deg(D) = 3$$

$$\deg(B) = \deg(E) = 4$$


$$\deg(F) = 1$$

$$\deg(G) = 0$$

- If G is a **directed** graph

In-degree: # of edges with the terminal vertex

Out-degree: # of edges with the initial vertex



$$\text{in_deg}(A) = 0 \quad \text{out_deg}(A) = 3$$


$$\text{in_deg}(B) = 1 \quad \text{out_deg}(B) = 3$$

$$\text{in_deg}(C) = 1 \quad \text{out_deg}(C) = 2$$

$$\text{in_deg}(D) = 2 \quad \text{out_deg}(D) = 1$$

Sources and Sinks

- For a **directed** graph,
 - Vertices with an **in-degree of zero** are described as **sources**
 - Vertices with an **out-degree of zero** are described as **sinks**




Sources: *A* and *G*


Sinks: *E*, *F*, and *G*

Subgraphs

A **subgraph** is a graph consisting of a **subset** of vertices and a **subset** of edges that connected those vertices in the original graph




the original graph




a subgraph

Subgraphs

A **subgraph** is a graph consisting of a **subset** of vertices and a **subset** of edges that connected those vertices in the original graph




the original graph




a subgraph

Vertex-Induced Subgraphs

A **vertex-induced subgraph** is a graph consisting of a subset of vertices where the edges are all edges in the original graph




the original graph




a vertex-induced subgraph

Vertex-Induced Subgraphs

A **vertex-induced subgraph** is a graph consisting of a subset of vertices where the edges are all edges in the original graph



the original graph




a vertex-induced subgraph

Paths

A **path** is an ordered sequence of vertices $\langle v_i, v_{i+1}, \dots, v_k \rangle$ where (v_j, v_{j+1}) is an edge for $j = i, \dots, k - 1$


- If G is an **undirected** graph
- If G is a **directed** graph



$\langle A, B, E, C, F \rangle$ is a path of length 4

$\langle A, B, E, C, B, D \rangle$ is a path of length 5

$\langle A \rangle$ is a path of length 0




$\langle A, B, D, E \rangle$ is a path of length 3


Connectivity

Two vertices v_i and v_j are said to be **connected** if there exists a path between them

A graph is **connected** if there exists a path between any two vertices




a connected graph



an unconnected graph

Strongly and Weakly Connectivity

- A graph is **strongly connected** if every pair of vertices (v_i and v_j) contains a path between **each other**
- A graph is **weakly connected** if there does not exist any path between any two pairs of vertices




Strongly connected subgraph:
 $\{C, D, E, H\}$

Weakly connected subgraph:
 $\{A, B, C, D, E, H\}$


Weighted Graphs

A **weighted graph** is a graph $G = (V, E, W)$ in which a weight $w_{ij} \in W$ is associated with each edge $(v_i, v_j) \in E$

- If G is an **undirected** graph
- If G is a **directed** graph




$\langle A, D, G \rangle$ is a path of length 8.8





$\langle A, D, B \rangle$ is a path of length 14.2

Trees

- Trees are **special cases** of a graph
 - Each is **connected**, and
 - There is a **unique path** between any two vertices



the original graph




example of trees t_1 and t_2


- **Consequences**
 - The number of edges is $|E| = |V| - 1$
 - The graph (tree) is **acyclic**
 - Adding one edge creates a cycle
 - Removing any one edge creates two disjoint non-empty subgraphs

Forests

- A **forest** is any graph that has no cycles
 - The number of edges is $|E| < |V|$
 - The number of trees is $|V| - |E|$
 - Removing any one edge adds one more tree to the forest



the original graph





example of a forest with 4 trees

Directed Acyclic Graphs

A **directed acyclic graph** (DAG) is a **directed** graph which has **no cycles**

- Graphical representation of **partial order** on a finite number of elements



example of two DAGs

Note that adding the edge (C, A) will make them not DAG

Graph Representations


- Space and time are analyzed in terms of
 - Number of vertices ($|V|$) and
 - Number of edges ($|E|$)
- There are two ways of representing graphs:
 - The adjacency matrix representation
 - The adjacency list representation

Adjacency Matrix

Require more memory but faster

- If G is an **undirected** graph

$$m_{ij} = \begin{cases} 1 & \text{if } (v_i, v_j) \in E \\ 0 & \text{otherwise} \end{cases}$$



$$M = \begin{matrix} & \begin{matrix} A & B & C & D & E & F \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \\ E \\ F \end{matrix} & \begin{pmatrix} 0 & 1 & 0 & 1 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \end{pmatrix}$$


- Memory requirement is $\Theta(|V|^2)$
- Determining if v_i is adjacent to v_j is $\Theta(1)$
- Finding all neighbors of v_i is $\Theta(|V|)$

Adjacency Matrix

Require more memory but faster

- If G is a **directed** graph

$$m_{ij} = \begin{cases} 1 & \text{if } (v_i, v_j) \in E \\ 0 & \text{otherwise} \end{cases}$$





$$M = \begin{matrix} & \begin{matrix} A & B & C & D & E & F \end{matrix} \\ \begin{matrix} A \\ B \\ C \\ D \\ E \\ F \end{matrix} & \begin{pmatrix} 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 \end{pmatrix} \end{pmatrix}$$

- Memory requirement is $\Theta(|V|^2)$
- Determining if v_i is adjacent to v_j is $\Theta(1)$
- Finding all neighbors of v_i is $\Theta(|V|)$

Adjacency List

Each vertex is associated with a list of its neighbors

- If G is an **undirected** graph





- Memory requirement is $\Theta(|V| + 2|E|) = \Theta(|V| + |E|)$
- On average, determining if v_i is adjacent to v_j is $\Theta(|E|/|V|)$
- On average, finding all neighbors of v_i is $\Theta(|E|/|V|)$

Adjacency List

Each vertex is associated with a list of its neighbors

- If G is a **directed** graph



- Memory requirement is $\Theta(|V| + |E|)$
- On average, determining if v_i is adjacent to v_j is $\Theta(|E|/|V|)$
- On average, finding all neighbors of v_i is $\Theta(|E|/|V|)$

Outline

- Graph ADT and Representations
- Graph Traversals
- Connectivity
- Bipartite Graphs
- Topological Sort

Strategies

- Traversals of graphs are also called **searches**
- We can use either
 - **Breadth-first traversal** requiring a queue, or
 - **Depth-first traversal** requiring a stack
- We will have to track which vertices have been visited, requiring $\Theta(|V|)$ memory

Breadth-First Traversal

Procedures:


- Choose a vertex, mark it as visited and insert into queue
- While the queue is not empty
 - Remove the front vertex v from the queue
 - For each vertex adjacent to v that has not been visited
 - Mark it visited, and
 - Insert it into the queue

Note that if the queue is empty and there are **no unvisited** vertices, the graph is **connected**

The size of the queue is $O(|V|)$

Example


Consider this graph:



Example: Breadth-First Traversal

Consider this graph:

- Start with the vertex *A*




Output: *A B C E D F G H I*

Time complexity: $\Theta(|V| + |E|)$

Example: Depth-First Traversal

Consider this graph:

- Start with the vertex *A*



Output: *A B C D E G I H F*

Time complexity: $\Theta(|V| + |E|)$

Applications

- Determining connectivity and finding connected subgraphs
- Determining the path length from one vertex to all others
- Testing if a graph is bipartite
- Determining maximum flow
- ...

Outline

- Graph ADT and Representations
- Graph Traversals
- Connectivity
- Bipartite Graphs
- Topological Sort


Connectivity

- Determine whether one vertex is **connected** to another
 - v_i is connected to v_j if there is a path from v_i to v_j
- **Strategy:**
 - Perform a **breadth-first traversal** starting at v_i
 - During the traversal,
 - If the vertex v_j ever found to be adjacent to any vertex in front of the queue, return true
 - Otherwise, if the queue is empty, return false

Determining Connectivity


Consider these graphs:

- Is A connected to D ?



BFS: $A B F G H E C D I$

↑
found here!



BFS: $A B F G E$


↑
NOT found

Connected Components

- Previously, if we continued the traversal, we would find all vertices that are connected to A
- Suppose we want to partition the vertices into **connected subgraphs**
 - While there are unvisited vertices in the graph
 - Select an unvisited vertex and perform a traversal on that vertex
 - Each vertex that is visited in that traversal is added to the set initially containing the initial unvisited vertex
 - Continue until all vertices are visited

Connected Components

Consider this graph:




<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>	<i>F</i>	<i>G</i>	<i>H</i>	<i>I</i>	<i>J</i>	<i>K</i>
1	2	6	8	9	11	5	3	4	7	10

There are three connected subgraphs:

$\{A, B, G, H, I\}$, $\{C, J\}$, and $\{D, E, F, K\}$

Strongly Connected Components

- A directed graph is **strongly connected** if there is a path between all pairs of vertices
- A **strongly connected component (SCC)** of a directed graph is a maximal strongly connected subgraph



Strongly Connected Components


Algorithm:

1. Call $\text{DFS}(G)$ to compute **exiting times** for each vertex
2. Compute G^T
3. Call $\text{DFS}(G^T)$, but in the main loop of DFS, consider the vertices in order of **decreasing times** as computed in (1)
4. Output the vertices of each tree in the depth-first forest of (3) as a separate strongly connected components

The algorithm takes $O(|V| + |E|)$ time


Strongly Connected Components

Step 1:



A	B	C	D	E	F	G	H	I	J
1	20	2	17	3	14	4	13	5	8
18	19	15	16	9	10	11	12	6	7

Steps 2-4:



A	B	C	D	E	F	G	H	I	J
1	8	3	4	9	14	11	12	17	20
5	6	2	7	10	13	15	16	18	19

Outline

- Graph ADT and Representations
- Graph Traversals
- Connectivity
- Bipartite Graphs
- Topological Sort

Definition



A **bipartite graph** is a graph where the vertices V can be divided into **two disjoint sets** V_1 and V_2 such that

- **Every edge** has one vertex in V_1 and the other in V_2
- In other words, no edges connects two vertices in the same set

Bipartite Graphs

Consider this graph:


- Is it a **bipartite** graph?
➡ **YES**, with a little work to decompose the vertices



Bipartite Graphs

Consider this graph:


- Is it a **bipartite** graph?
➔ **YES**, with a little work to decompose the vertices



Bipartite Graphs

Consider this graph:

- Is it a **bipartite** graph?
 - ➔ **NO**, we cannot divide $\{A, B, F\}$, for instance



Bipartite Graphs


Algorithm: using a breadth-first traversal

- Choose a vertex, mark it belonging to V_1 and insert in into a queue
- While the queue is not empty
 - Dequeue the front vertex v , and
 - Any adjacent vertices that are already marked must belong to the set not containing v , otherwise, the graph is not bipartite (we are done); while
 - Any unmarked adjacent vertices are marked as belonging to the other set and they are inserted into the queue
- If the queue is empty, the graph is bipartite

Bipartite Graphs

Consider the previous graph:

- We will use colors to distinguish the two sets




BFS: *A B F G H E C I D*

OK, this graph is bipartite

Bipartite Graphs

Consider the previous graph:

- We will use colors to distinguish the two sets




BFS: *A B F H E G C*


This graph is NOT bipartite

Complete Bipartite Graphs $K_{m,n}$


- Vertex set partitioned into two subsets of size m and n
- All vertices in one subset are connected to all vertices in the other subset



$K_{1,5}$



$K_{2,3}$



$K_{3,3}$

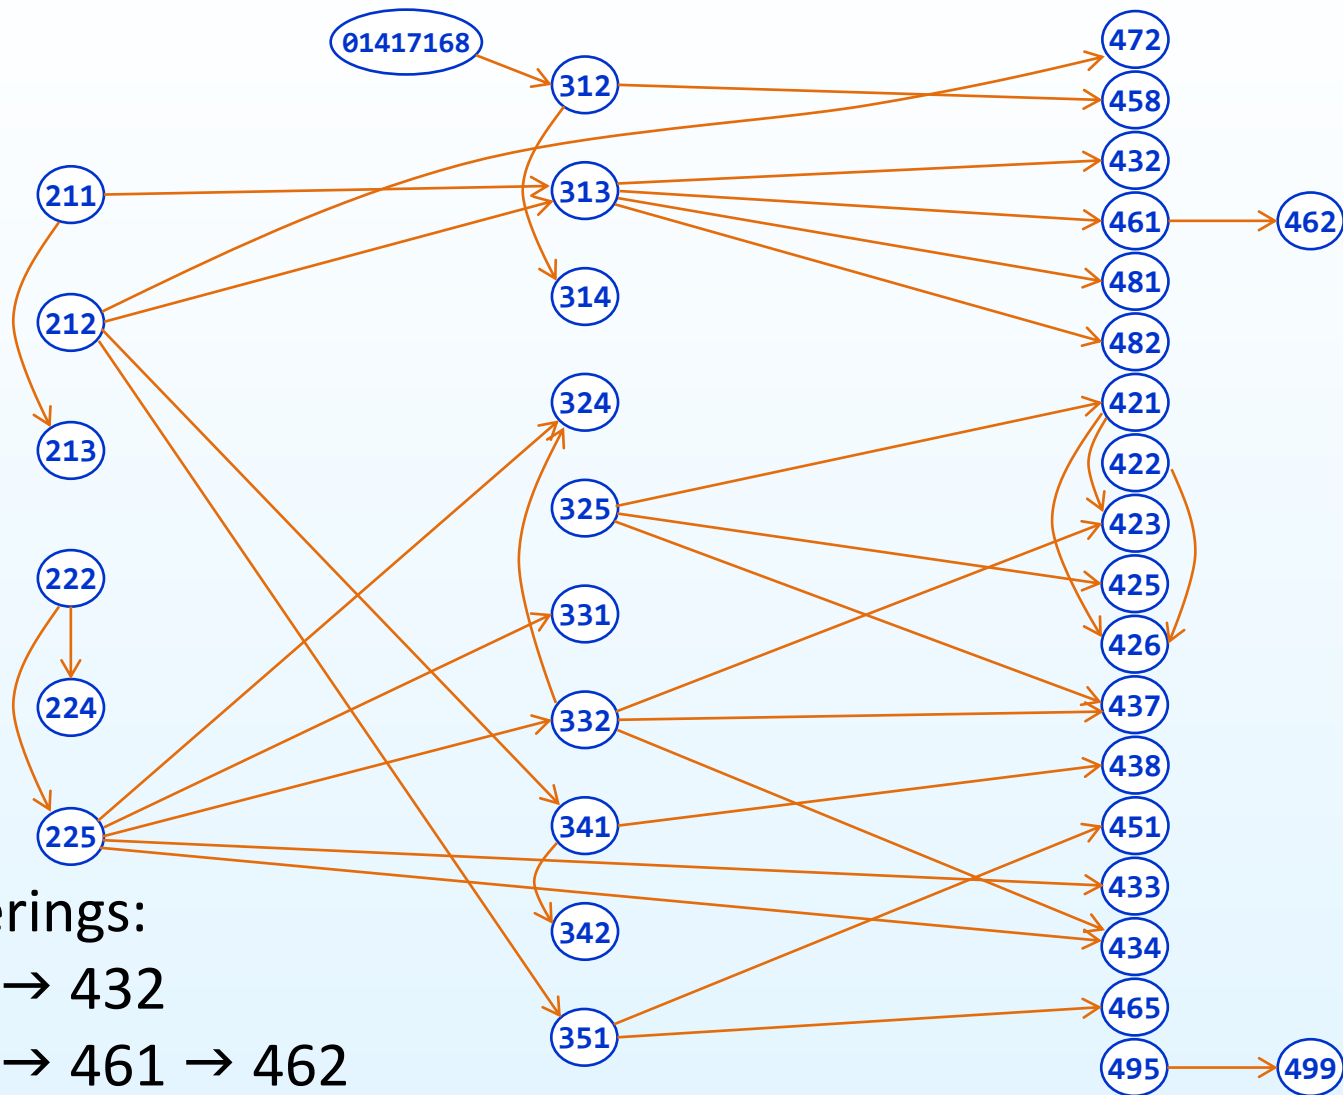
Outline

- Graph ADT and Representations
- Graph Traversals
- Connectivity
- Bipartite Graphs
- Topological Sort

Motivation

- Given a set of tasks with dependencies
 - Is there an order in which we can complete the tasks?
- Dependencies form a partial ordering
 - A partial ordering on a finite number of objects can be represented as a **directed acyclic graph** (DAG)

Example: CPE Course Prerequisites



Partial orderings:

212 → 313 → 432

212 → 313 → 461 → 462

Definition

Restriction of paths in DAGs

- Given two different vertices v_i and v_j , there **cannot both** be a path from v_i to v_j and a path from v_j to v_i

Topological sort

- A **topological sort** of the vertices in a DAG is an ordering

$$v_1, v_2, v_3, \dots, v_n$$


such that if there is a path from v_i to v_j then v_i appears before v_j

Example

Consider this DAG:

- A topological sort is

$C, H, D, I, A, J, B, F, G, K, E, L$




Note that

- There are paths from C , D , H , I , and J to F , so all these must come before F in a topological sort
- Clearly, this sorting need **not be unique**

Applications

Suppose you want to dress up to dinner

- You must wear the following:
 - jacket, shirt, briefs, pants, socks, tie, etc.
- There are certain constraints:
 - The pants should go on after the briefs
 - Socks are put on before shoes
- One topological sort is



briefs, pants, wallet, belt, keys, socks, shoes, shirt,
tie, jacket, phone, watch

- A more reasonable topological sort is
briefs, socks, pants, shirt, belt, keys, tie, jacket,
wallet, phone, watch, shoes

Topological Sort

Algorithm:


Given a graph, iterate the followings $|V|$ times

- Choose a vertex v that has in-degree zero
- Let v be the next vertex in our topological sort
- Remove v and all edges connected to it

Example

Consider the previous DAG:

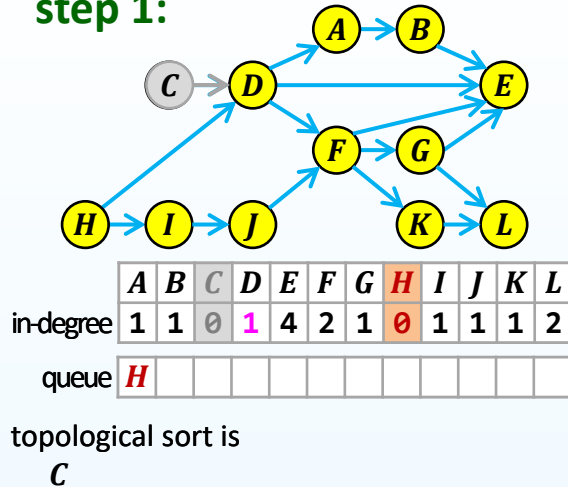
- First, create a table of in-degrees of each vertex
- Then, create a queue



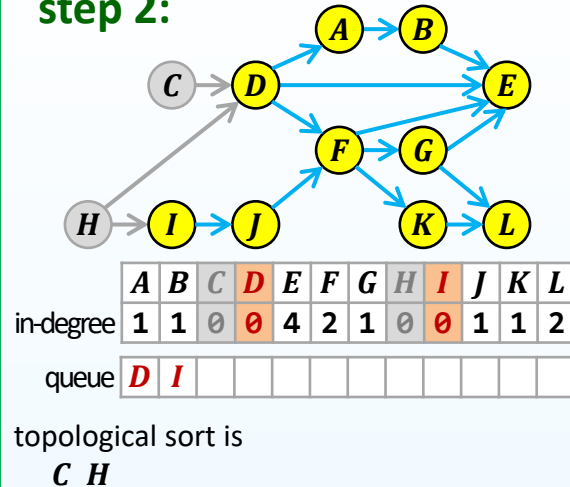
	<i>A</i>	<i>B</i>	<i>C</i>	<i>D</i>	<i>E</i>	<i>F</i>	<i>G</i>	<i>H</i>	<i>I</i>	<i>J</i>	<i>K</i>	<i>L</i>
in-degree	1	1	0	2	4	2	1	0	1	1	1	2
queue	<i>C</i>	<i>H</i>										

Example


step 1:



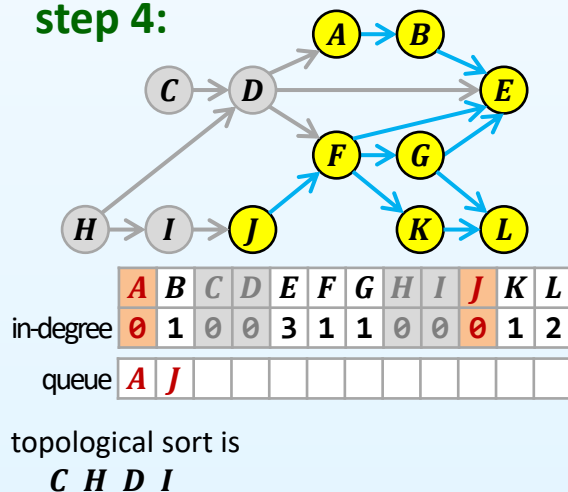
step 2:



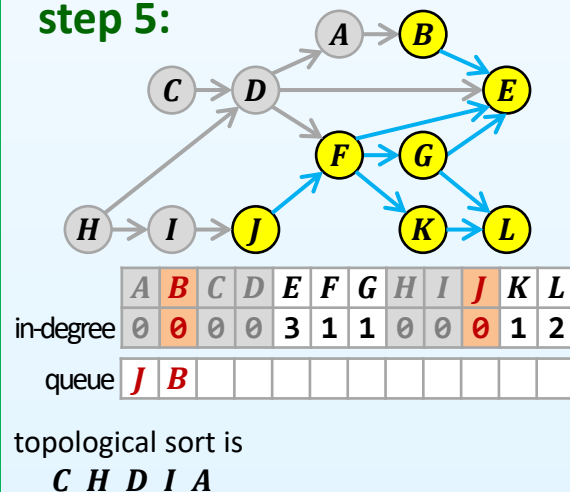
step 3:




step 4:



step 5:




step 6:




Example

step 7:




topological sort is
C H D I A J B

step 8:




topological sort is
C H D I A J B F

step 9:




topological sort is
C H D I A J B F G

step 10:




topological sort is
C H D I A J B F G K

step 11:



topological sort is
C H D I A J B F G K E

step 12:



topological sort is
C H D I A J B F G K E L

Complexity Analysis

- Initialize in-degree array
 - ➔ $\Theta(|E|)$
- Initialize queue with zero in-degree vertices
 - ➔ $\Theta(|V|)$
- Dequeue and output vertex
 - All vertices will be enqueued and dequeued once, each take $\Theta(1)$
 - ➔ $\Theta(|V|)$
- Reduce in-degree of all vertices adjacent to a vertex and enqueue any zero in-degree vertices
 - All edges will decrease in-degree of adjacent once, each take $\Theta(1)$
 - ➔ $\Theta(|E|)$

Finally, the topological sort takes $\Theta(|V| + |E|)$ time

Any Question?

